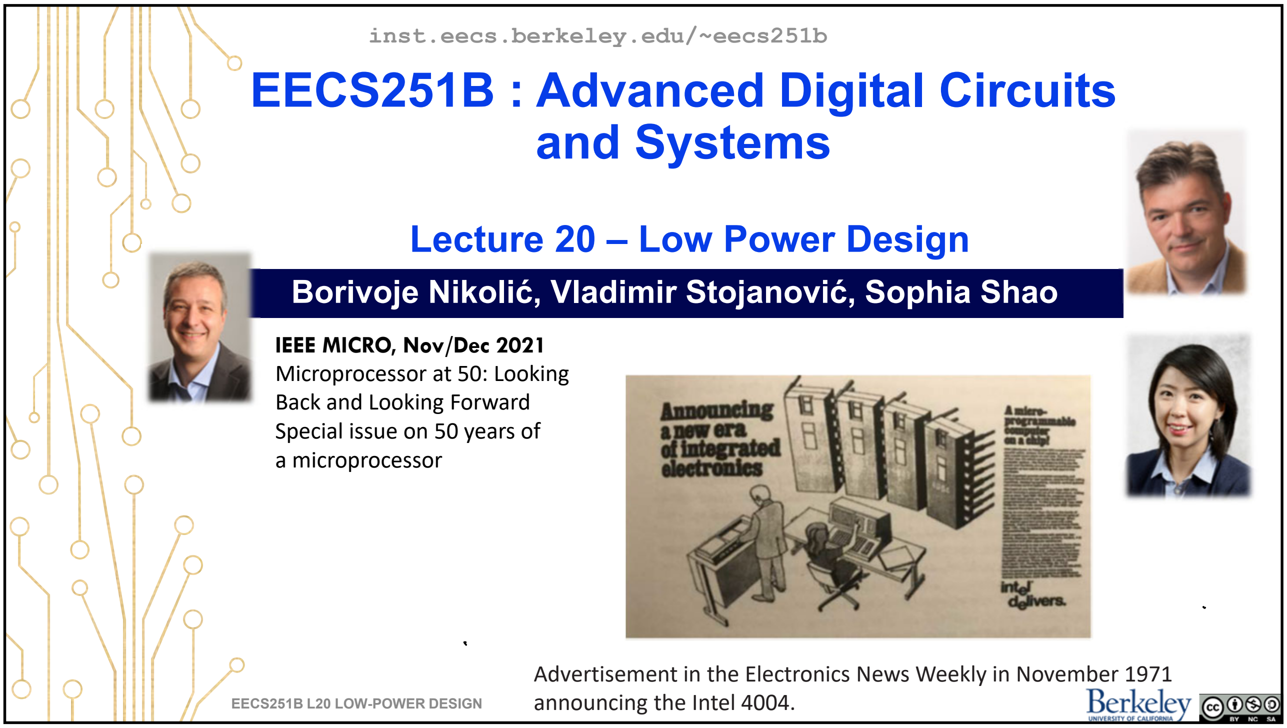

inst.eecs.berkeley.edu/~eecs251b

EECS251B : Advanced Digital Circuits and Systems

Lecture 20 – Low Power Design

Borivoje Nikolić, Vladimir Stojanović, Sophia Shao

IEEE MICRO, Nov/Dec 2021
Microprocessor at 50: Looking Back and Looking Forward
Special issue on 50 years of a microprocessor

Announcing a new era of integrated electronics
A micro-programmable computer on a chip!
intel delivers.

Advertisement in the Electronics News Weekly in November 1971 announcing the Intel 4004.

EECS251B L20 LOW-POWER DESIGN

Berkeley UNIVERSITY OF CALIFORNIA

1

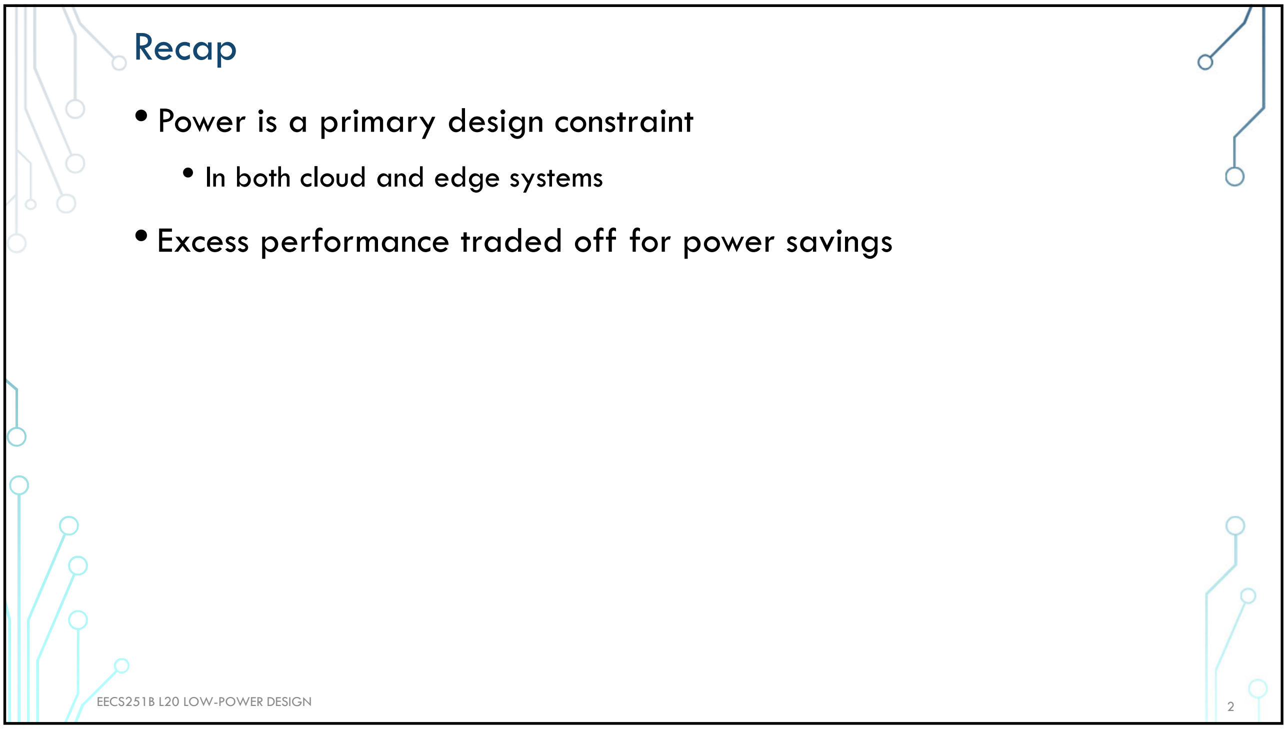

Recap

- Power is a primary design constraint
  - In both cloud and edge systems
- Excess performance traded off for power savings

EECS251B L20 LOW-POWER DESIGN

2

2

Architectural Optimizations

Optimal Processors

- Processors used to be optimized for performance
  - Optimal logic depth was found to be 8-11 FO4 delays in superscalar processors
  - 1.8-3 FO4 in sequentials, rest in combinatorial
    - Kunkel, Smith, ISCA'86
    - Hriskesh, Jouppi, Farkas, Burger, Keckler, Shivakumar, ISCA'02
    - Harstein, Puzak, ISCA'02
    - Sprangle, Carmean, ISCA'02
- But those designs have very high power dissipation
  - Need to optimize for both performance and power/energy

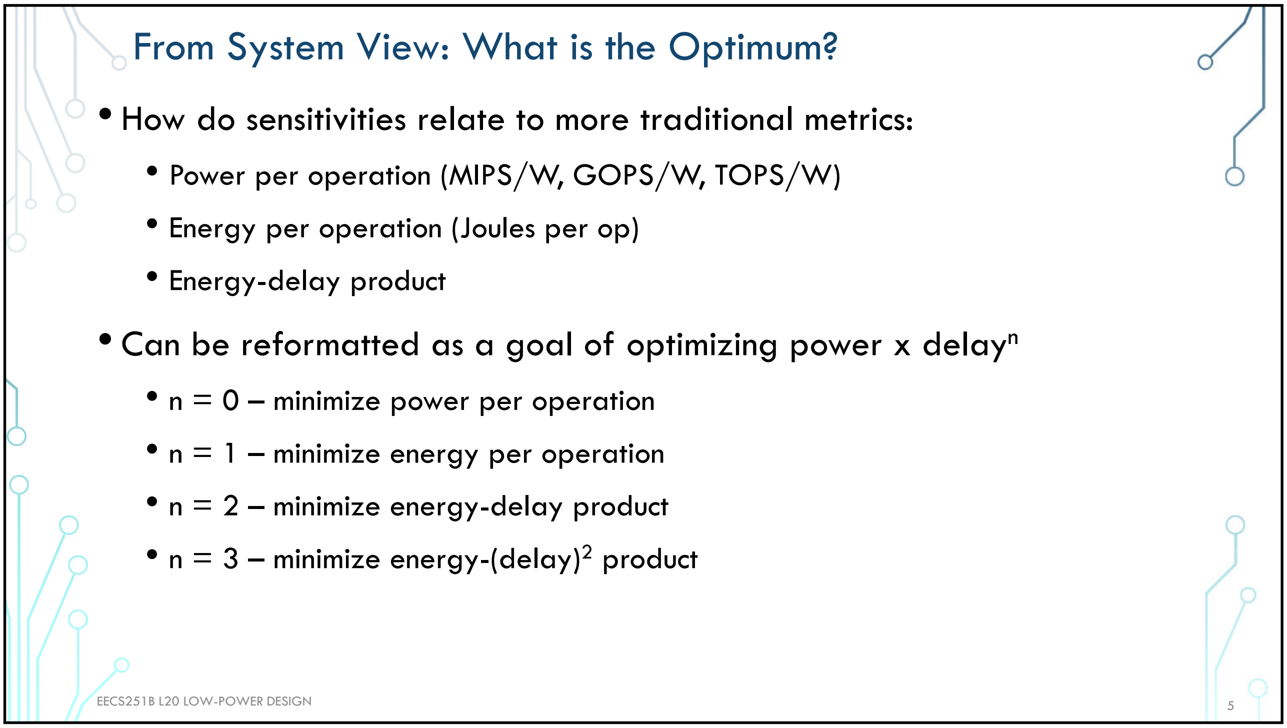

From System View: What is the Optimum?

- How do sensitivities relate to more traditional metrics:
  - Power per operation (MIPS/W, GOPS/W, TOPS/W)
  - Energy per operation (Joules per op)
  - Energy-delay product
- Can be reformatted as a goal of optimizing power x delay^n
  - n = 0 – minimize power per operation
  - n = 1 – minimize energy per operation
  - n = 2 – minimize energy-delay product
  - n = 3 – minimize energy-(delay)^2 product

5

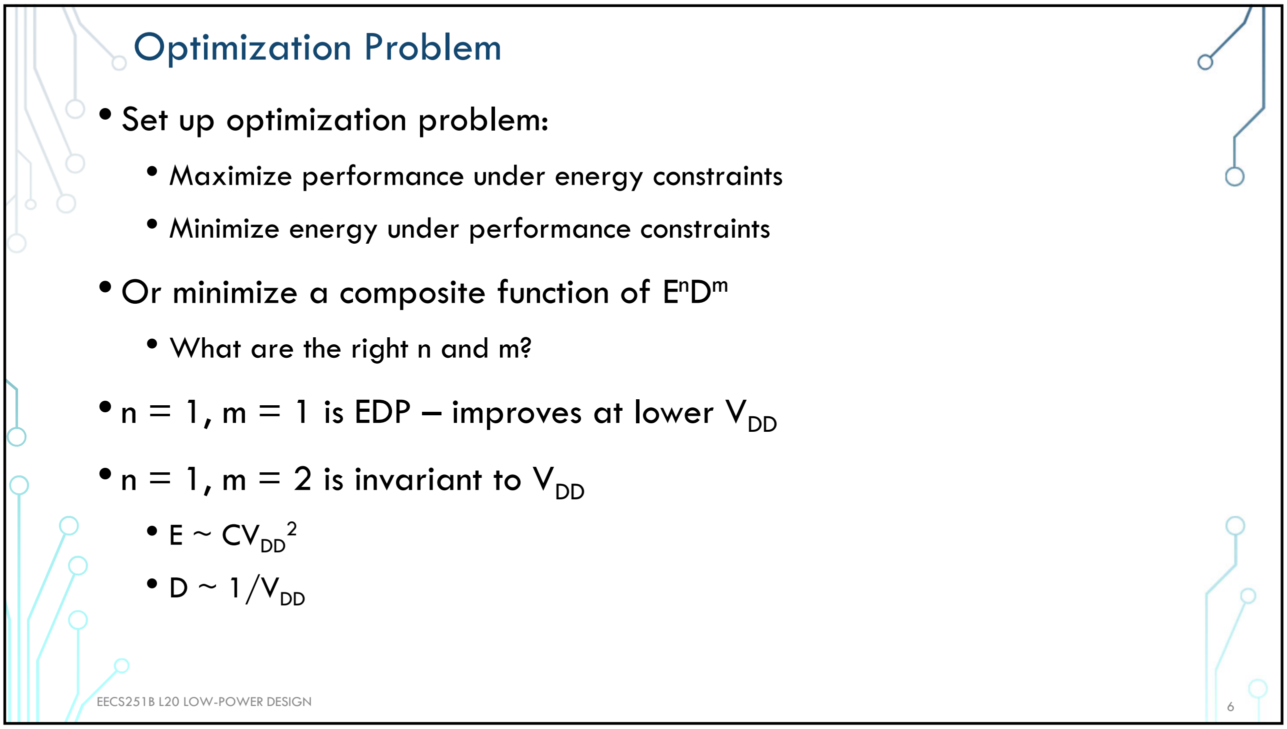

Optimization Problem

- Set up optimization problem:
  - Maximize performance under energy constraints
  - Minimize energy under performance constraints
- Or minimize a composite function of E^nD^m
  - What are the right n and m?
- n = 1, m = 1 is EDP – improves at lower V_{DD}
- n = 1, m = 2 is invariant to V_{DD}
  - $E \sim CV_{DD}^2$
  - $D \sim 1/V_{DD}$

6

Hardware Intesnity

- Introduced by Zyuban and Strenski in 2002.
- Measures where is the design on the Energy-Delay curve
- Parameter in cost function optimization

$$F_c = (E/E_0)(D/D_0)^{\eta} \qquad 0 \leq \eta < +\infty,$$

$$\eta = -\left.\frac{D \partial E}{E \partial D}\right|_{v}$$

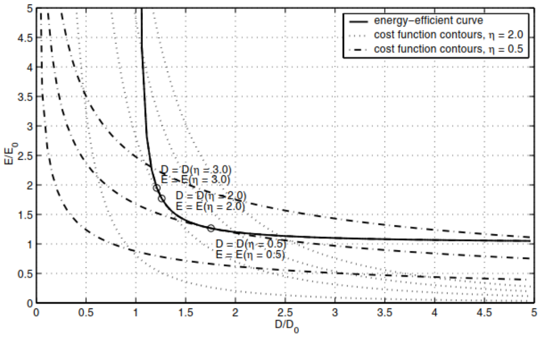

Slope of the optimal E-D curve at the chosen design point

Optimum Across Hierarchy Layers

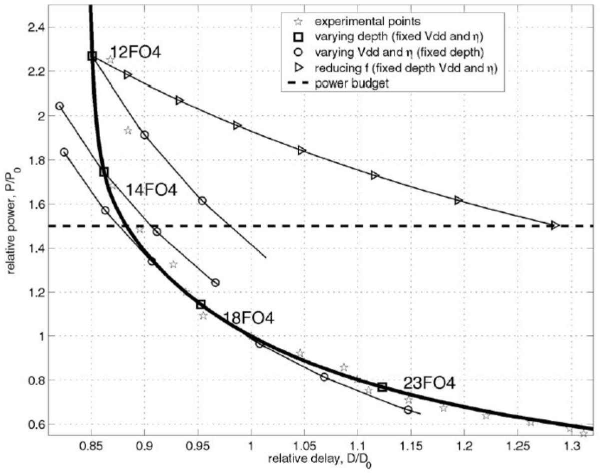

Zyuban et al, TComp'04

Optimal logic depth in pipelined processors is ~18FO4

Relatively flat in the 16-22FO4 range

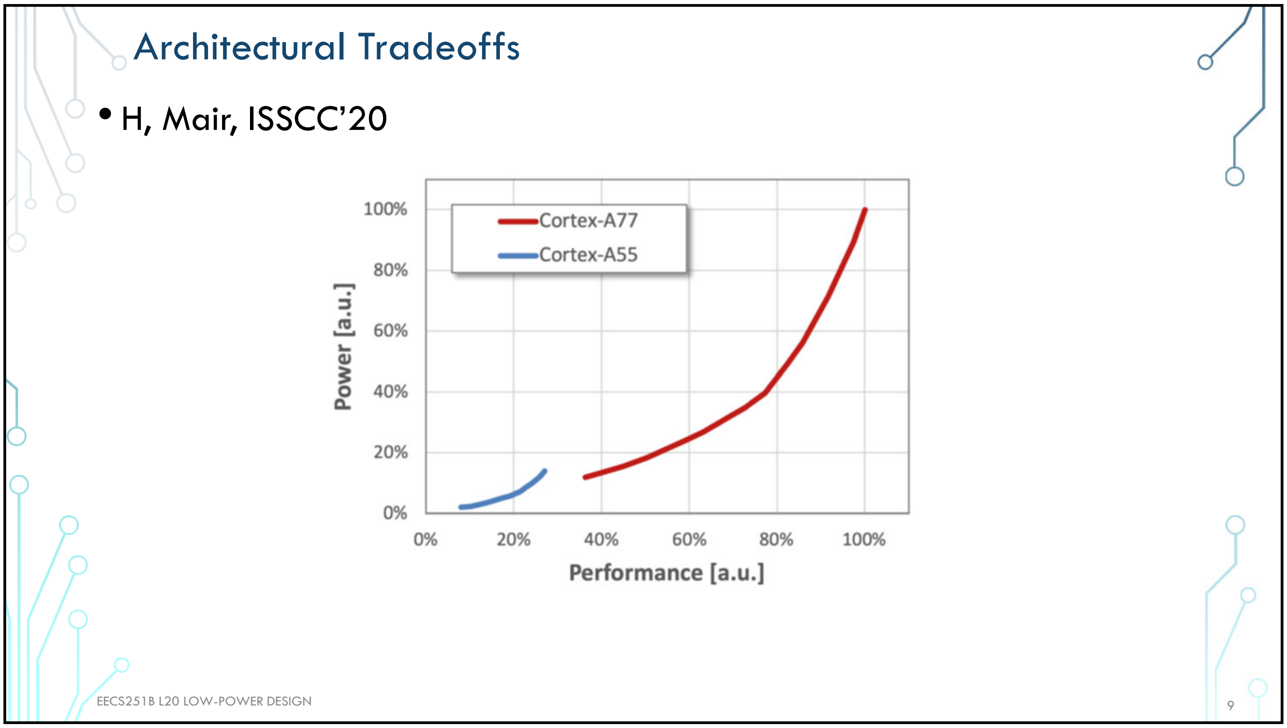

Architectural Tradeoffs

- H, Mair, ISSCC’20

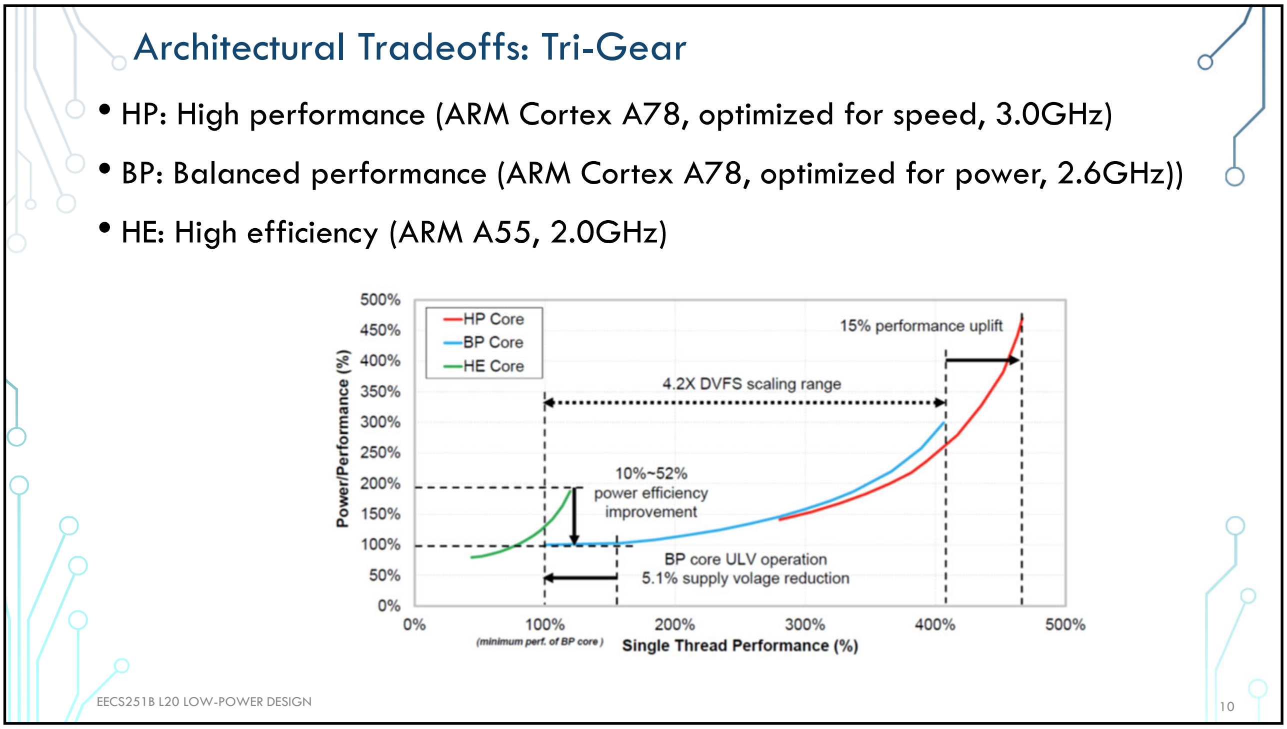

Architectural Tradeoffs: Tri-Gear

- HP: High performance (ARM Cortex A78, optimized for speed, 3.0GHz)
- BP: Balanced performance (ARM Cortex A78, optimized for power, 2.6GHz))
- HE: High efficiency (ARM A55, 2.0GHz)

Announcements

- Quiz 2 today
- Homework 3 due next week

11

Circuit-Level Tradeoffs

12

Alpha-Power Based Delay Model

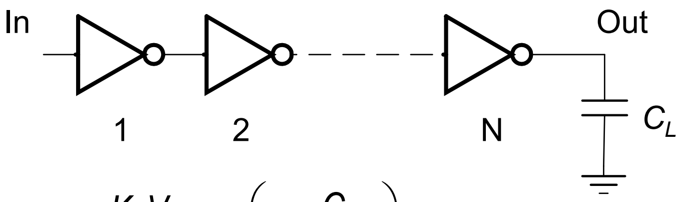

$$t_{pi} = \frac{K_d V_{DD}}{(V_{DD}-V_{Th})^{\alpha}}\left(1+\frac{C_{L,i}}{C_{in,i}}\right)$$

$$D = \sum t_{pi} = \sum \frac{K_d V_{DD}}{(V_{DD}-V_{Th})^{\alpha}}\left(1+\frac{W_{L,i}}{W_{in,i}}\right)$$

Energy Models

- **Switching**

$$E_{Sw} = \alpha_{0\to 1}\left(C_{L,i}+C_{\text{int},i}\right)V_{DD}{}^{2}$$

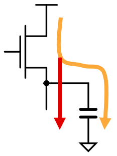

- **Leakage**

$$E_{Lk} = W_{In} I_0 e^{\frac{-(V_{Th}-\gamma V_{DD})}{nV_t}} V_{DD} D$$

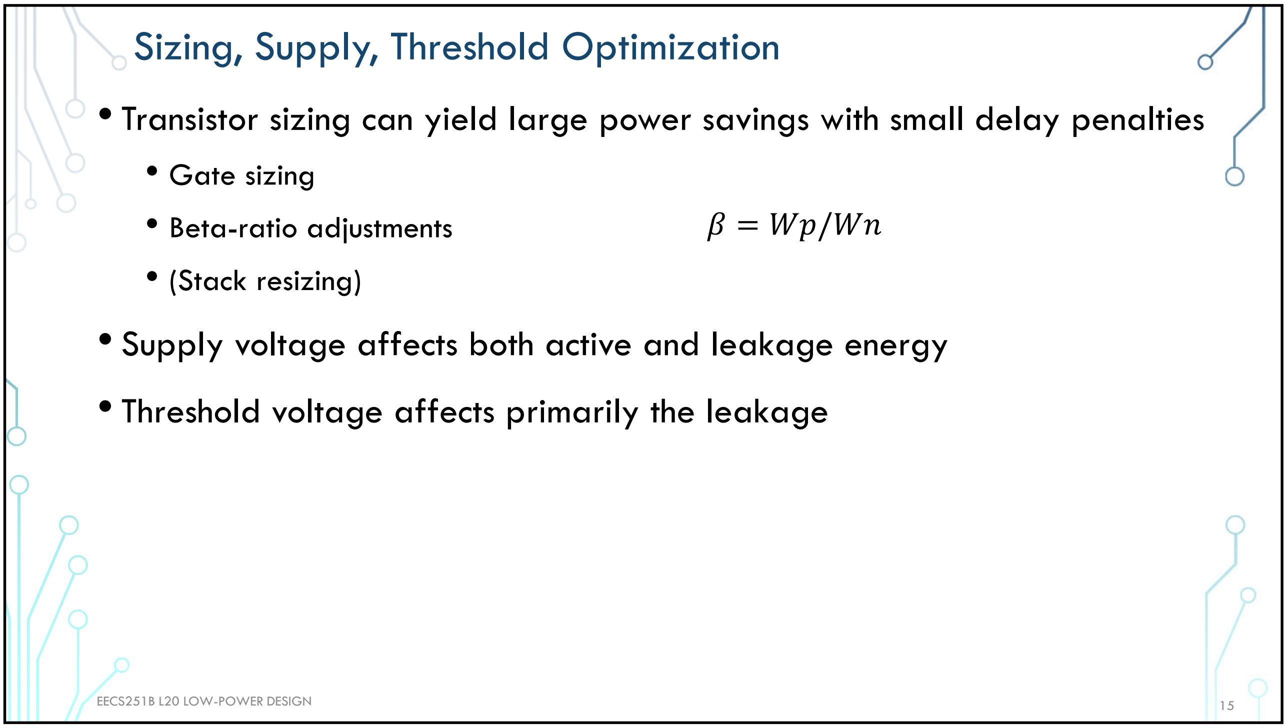

Sizing, Supply, Threshold Optimization

- Transistor sizing can yield large power savings with small delay penalties
  - Gate sizing
  - Beta-ratio adjustments $\beta = Wp/Wn$
  - (Stack resizing)
- Supply voltage affects both active and leakage energy
- Threshold voltage affects primarily the leakage

15

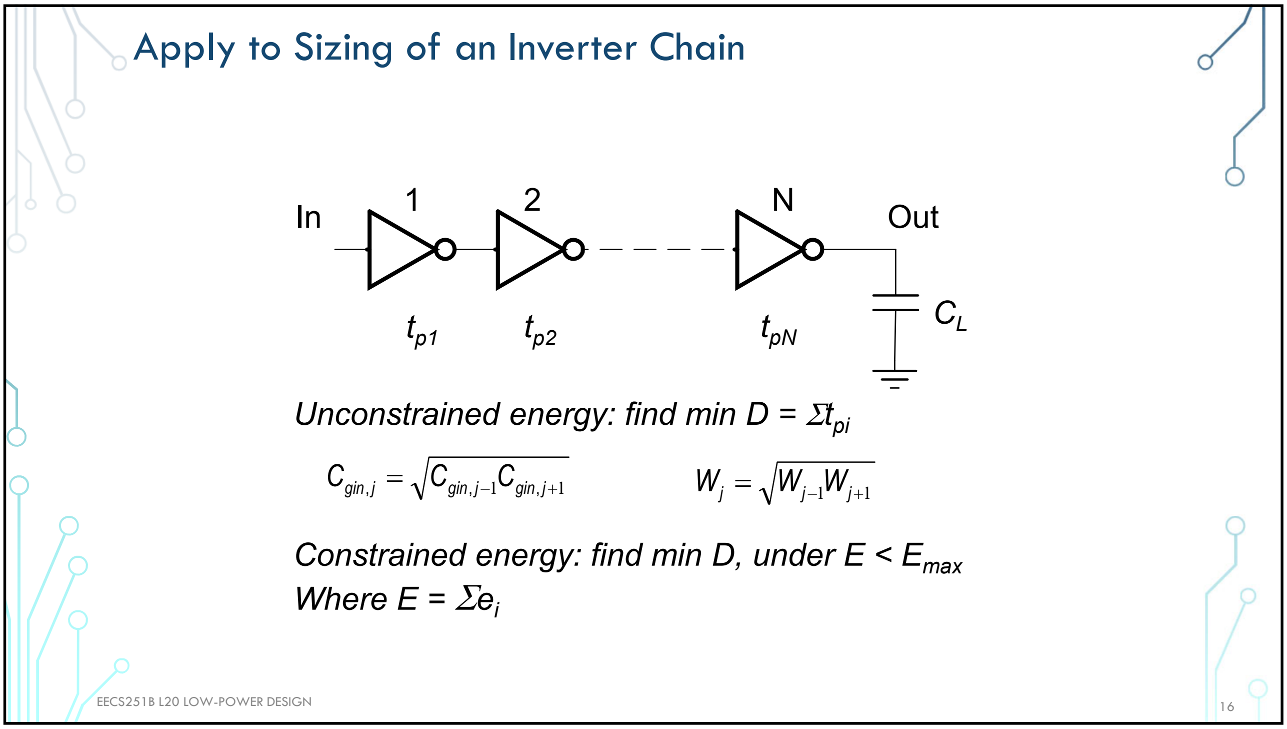

Apply to Sizing of an Inverter Chain

Unconstrained energy: find min $D = \Sigma t_{pi}$

$$C_{gin,j} = \sqrt{C_{gin,j-1} C_{gin,j+1}} \qquad W_j = \sqrt{W_{j-1} W_{j+1}}$$

Constrained energy: find min D, under $E < E_{max}$
Where $E = \Sigma e_i$

16

Constrained Optimization

- Find min(D) subject to $E = E_{max}$
  - *Constrained function minimization*
- E.g. Lagrange multipliers

$$\Lambda(x) = D(x) + \lambda\left(E(x) - E_{\text{max}}\right)$$

$$\frac{\partial \Lambda}{\partial x} = 0$$

Or dual:

$$\mathrm{K}(x) = E(x) + \lambda\left(D - D_{\text{max}}\right)$$

- Can solve analytically for $x = W_i$, V_{DD}, V_{Th}

Inverter Chain: Sizing Optimization

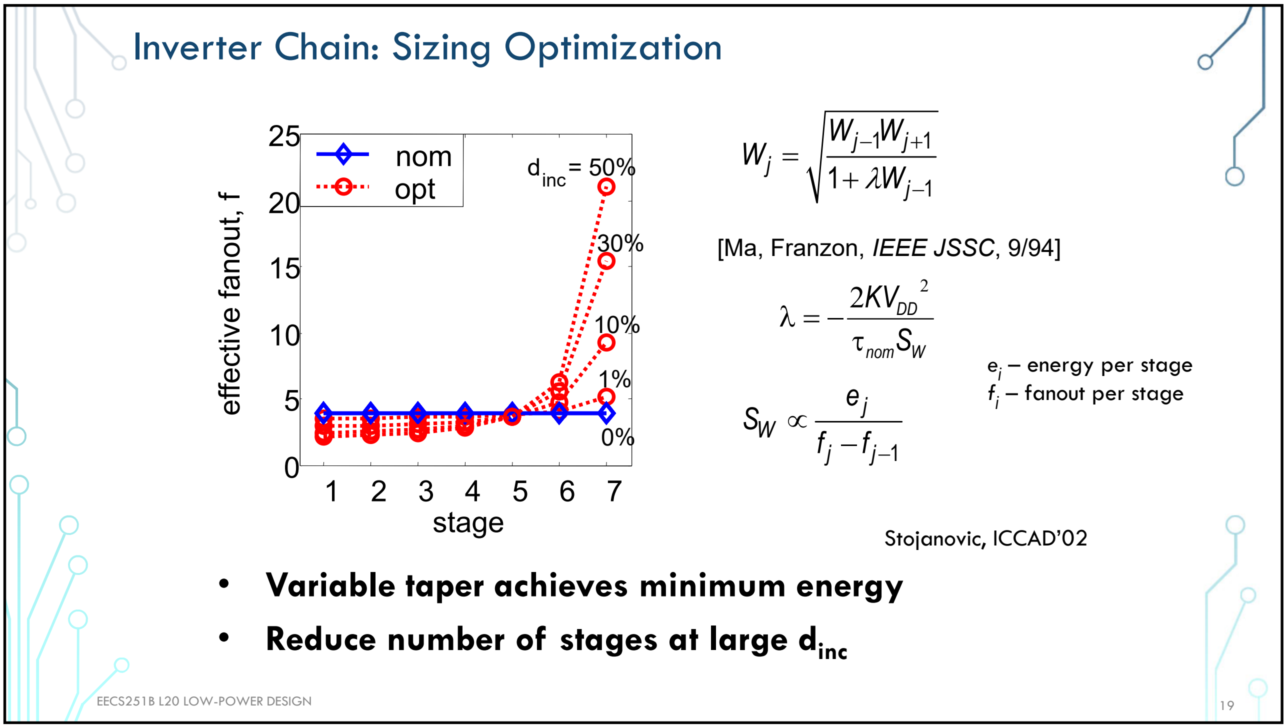

Inverter Chain: Sizing Optimization

$$W_j = \sqrt{\frac{W_{j-1}W_{j+1}}{1+\lambda W_{j-1}}}$$

[Ma, Franzon, *IEEE JSSC*, 9/94]

$$\lambda = -\frac{2KV_{DD}^{2}}{\tau_{nom}S_W}$$

$$S_W \propto \frac{e_j}{f_j - f_{j-1}}$$

e_j – energy per stage
f_j – fanout per stage

Stojanovic, ICCAD'02

- **Variable taper achieves minimum energy**
- **Reduce number of stages at large d_{inc}**

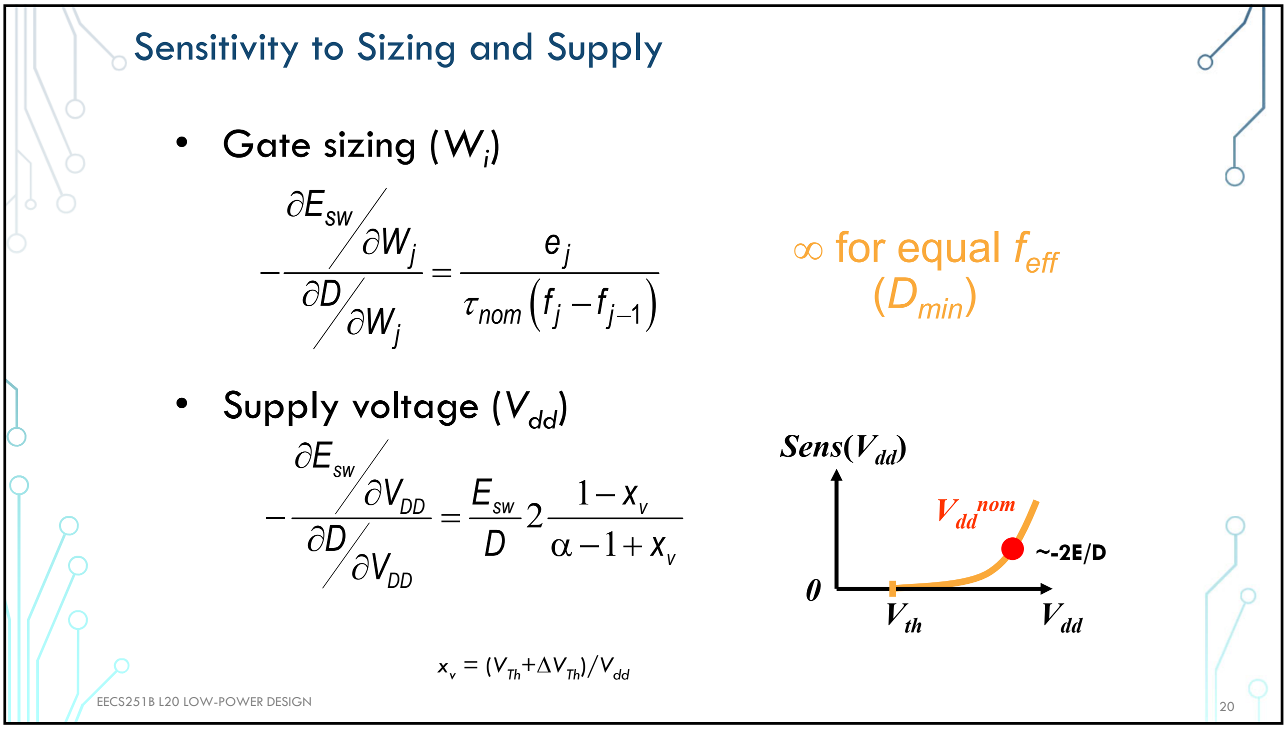

Sensitivity to Sizing and Supply

- Gate sizing (W_i)

$$-\frac{\partial E_{sw}/\partial W_j}{\partial D/\partial W_j} = \frac{e_j}{\tau_{nom}\left(f_j - f_{j-1}\right)}$$

∞ for equal f_{eff}
(D_{min})

- Supply voltage (V_{dd})

$$-\frac{\partial E_{sw}/\partial V_{DD}}{\partial D/\partial V_{DD}} = \frac{E_{sw}}{D}2\frac{1-x_v}{\alpha-1+x_v}$$

$x_v = (V_{Th}+\Delta V_{Th})/V_{dd}$

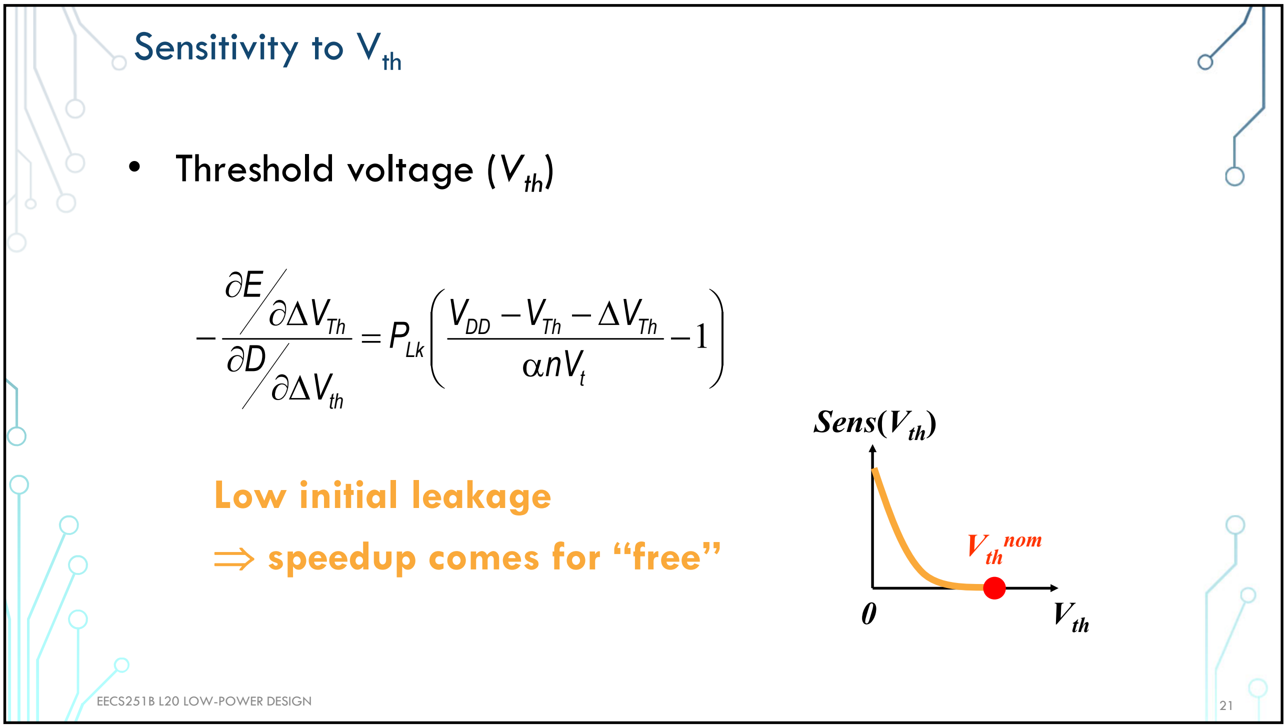

Sensitivity to V_{th}

- Threshold voltage (V_{th})

$$-\frac{\partial E/\partial \Delta V_{Th}}{\partial D/\partial \Delta V_{th}} = P_{Lk}\left(\frac{V_{DD}-V_{Th}-\Delta V_{Th}}{\alpha n V_t}-1\right)$$

Low initial leakage

⇒ speedup comes for "free"

$Sens(V_{th})$
V_{th}^{nom}
0
V_{th}

Scaling Supplies

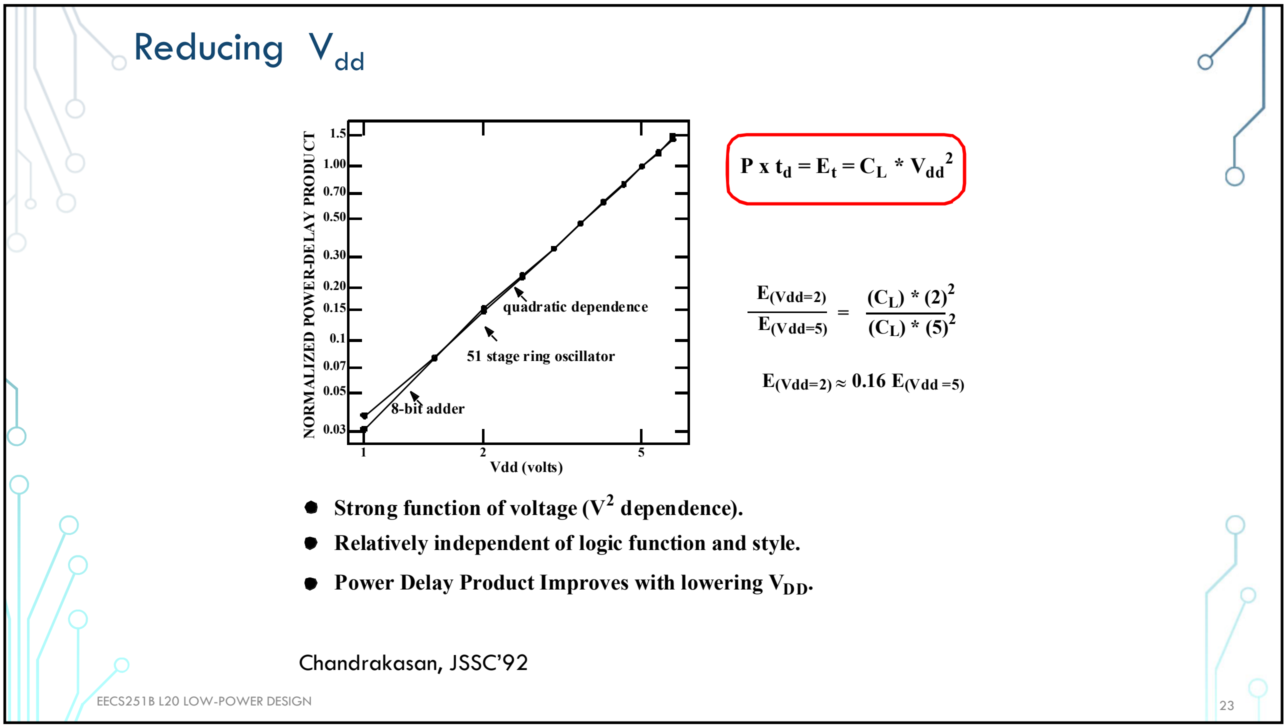

Reducing V_{dd}

$$P \times t_d = E_t = C_L * V_{dd}^2$$

$$\frac{E_{(Vdd=2)}}{E_{(Vdd=5)}} = \frac{(C_L) * (2)^2}{(C_L) * (5)^2}$$

$$E_{(Vdd=2)} \approx 0.16\ E_{(Vdd=5)}$$

- Strong function of voltage (V^2 dependence).
- Relatively independent of logic function and style.
- Power Delay Product Improves with lowering V_{DD}.

Chandrakasan, JSSC'92

23

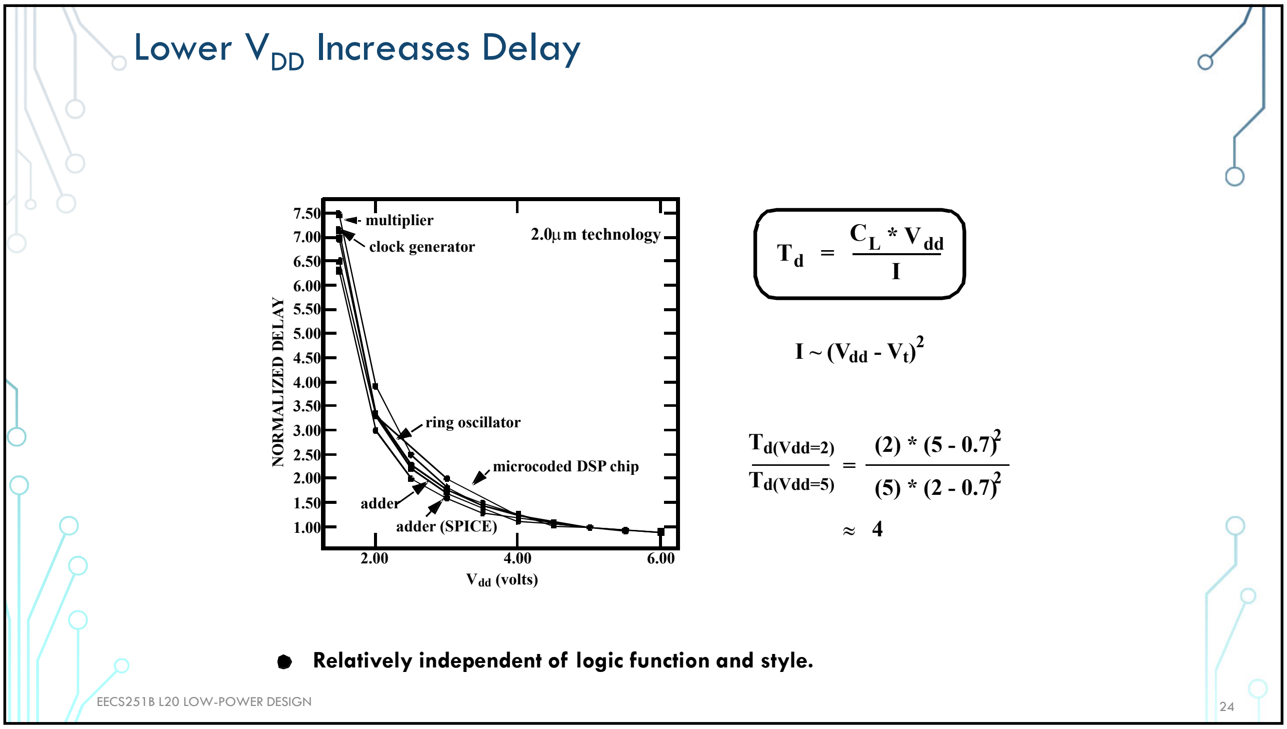

Lower V_{DD} Increases Delay

$$T_d = \frac{C_L * V_{dd}}{I}$$

$$I \sim (V_{dd} - V_t)^2$$

$$\frac{T_{d(Vdd=2)}}{T_{d(Vdd=5)}} = \frac{(2) * (5 - 0.7)^2}{(5) * (2 - 0.7)^2}$$

$$\approx 4$$

- Relatively independent of logic function and style.

24

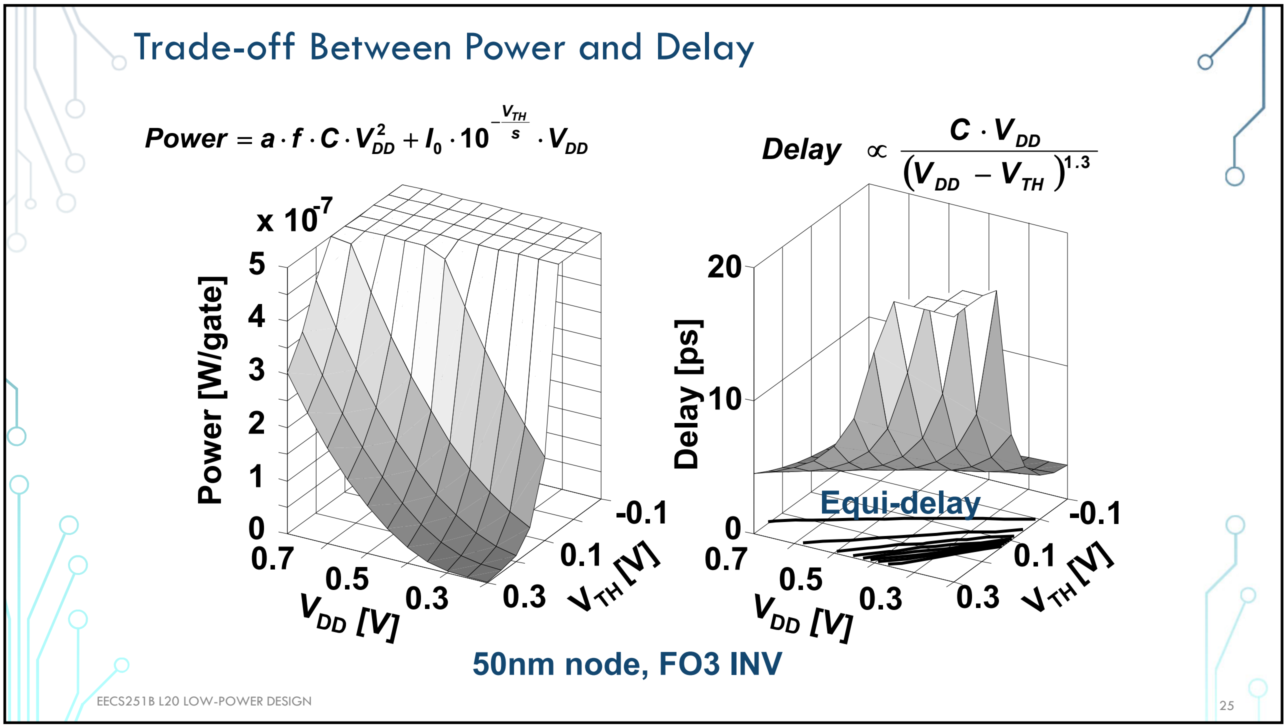

Trade-off Between Power and Delay

$$Power = a \cdot f \cdot C \cdot V_{DD}^2 + I_0 \cdot 10^{-\frac{V_{TH}}{s}} \cdot V_{DD}$$

$$Delay \propto \frac{C \cdot V_{DD}}{(V_{DD} - V_{TH})^{1.3}}$$

50nm node, FO3 INV

EECS251B L20 LOW-POWER DESIGN

25

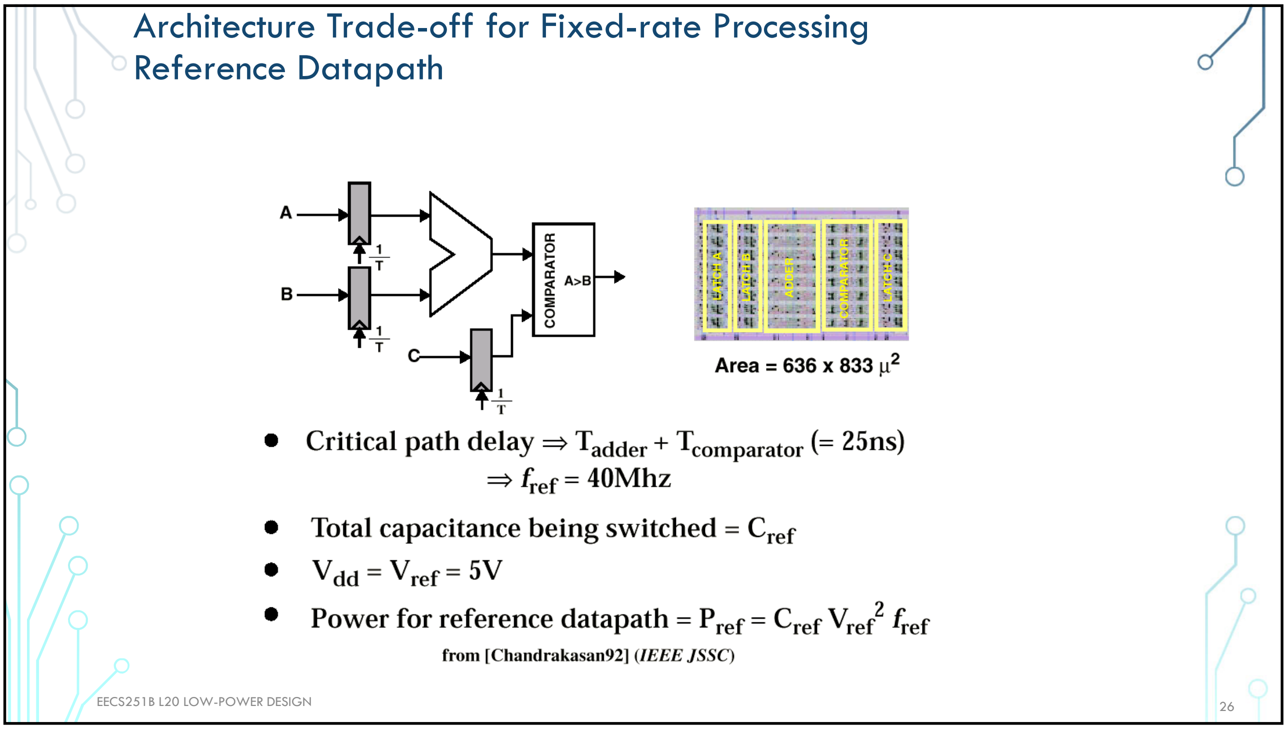

Architecture Trade-off for Fixed-rate Processing Reference Datapath

Area = 636 x 833 μ^2

- Critical path delay ⇒ $T_{adder} + T_{comparator}$ (= 25ns)
  ⇒ f_{ref} = 40Mhz
- Total capacitance being switched = C_{ref}
- $V_{dd} = V_{ref} = 5V$
- Power for reference datapath = $P_{ref} = C_{ref} V_{ref}^2 f_{ref}$

from [Chandrakasan92] (*IEEE JSSC*)

EECS251B L20 LOW-POWER DESIGN

26

Parallel Datapath

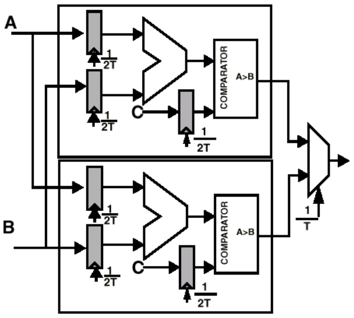

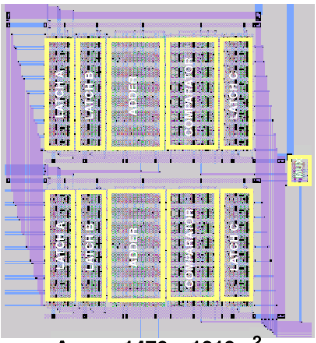

Area = 1476 x 1219 μ^2

- The clock rate can be reduced by half with the same throughput ⇒ $f_{par} = f_{ref} / 2$
- $V_{par} = V_{ref} / 1.7$, $C_{par} = 2.15C_{ref}$
- $P_{par} = (2.15C_{ref})\ (V_{ref}/1.7)^2\ (f_{ref}/2) \approx 0.36\ P_{ref}$

Pipelined Datapath

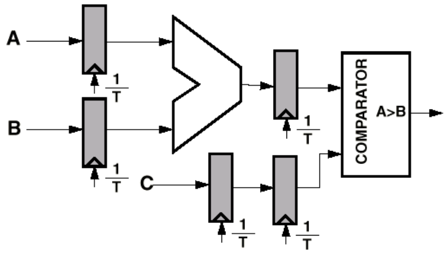

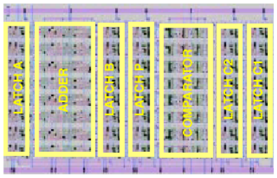

Area = 640 x 1081 μ^2

- Critical path delay is less ⇒ max $[T_{adder}, T_{comparator}]$
- Keeping clock rate constant: $f_{pipe} = f_{ref}$
  Voltage can be dropped ⇒ $V_{pipe} = V_{ref} / 1.7$
- Capacitance slightly higher: $C_{pipe} = 1.15C_{ref}$
- $P_{pipe} = (1.15C_{ref})\ (V_{ref}/1.7)^2\ f_{ref} \approx 0.39\ P_{ref}$

A Simple Datapath: Summary

Architecture type	Voltage	Area	Power
Simple datapath (no pipelining or parallelism)	5V	1	1
Pipelined datapath	2.9V	1.3	0.39
Parallel datapath	2.9V	3.4	0.36
Pipeline-Parallel	2.0V	3.7	0.2

Multiple Supplies

Multiple Supply Voltages

- Block-level supply assignment (“power domains” or “voltage islands”)
  - Higher throughput/lower latency functions are implemented in higher V_{DD}
  - Slower functions are implemented with lower V_{DD}
  - Often called “Voltage islands”
  - Separate supply grids, level conversion performed at block boundaries
- Multiple supplies inside a block
  - Non-critical paths moved to lower supply voltage
  - Level conversion within the block
  - Physical design challenging
  - (Not used in practice)

Power Domains

Practical Examples

- Intel Skylake (ISSCC'16)
  - Four power planes indicated by colors

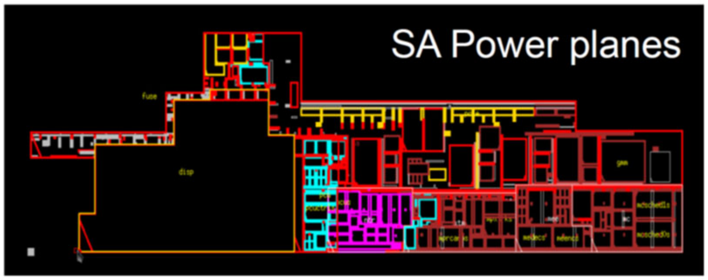

Practical Examples

- Intel 28-core Skylake-SP (ISSCC'18)

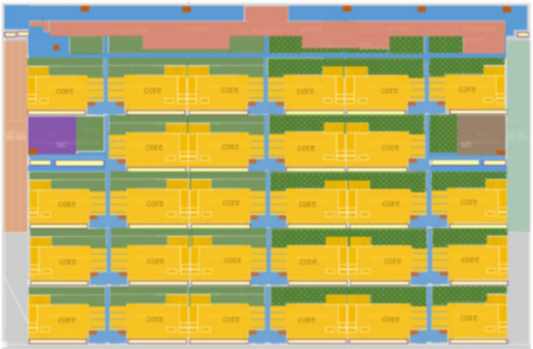

🟨 Vcc: core supply (per core)

🟩🟩 } Vccclm: Un-core supply

🟥 Vccsa: System Agent supply

🟦 Vccio: Infrastructure supply

🟥 Vccsfr: PLL supply

🟪🟫 } Vccddrd: DDR logic supply

🟧🟩 } Vccddra: DDR I/O supply

- 9 primary VCC domains are partitioned into 35 VCC planes

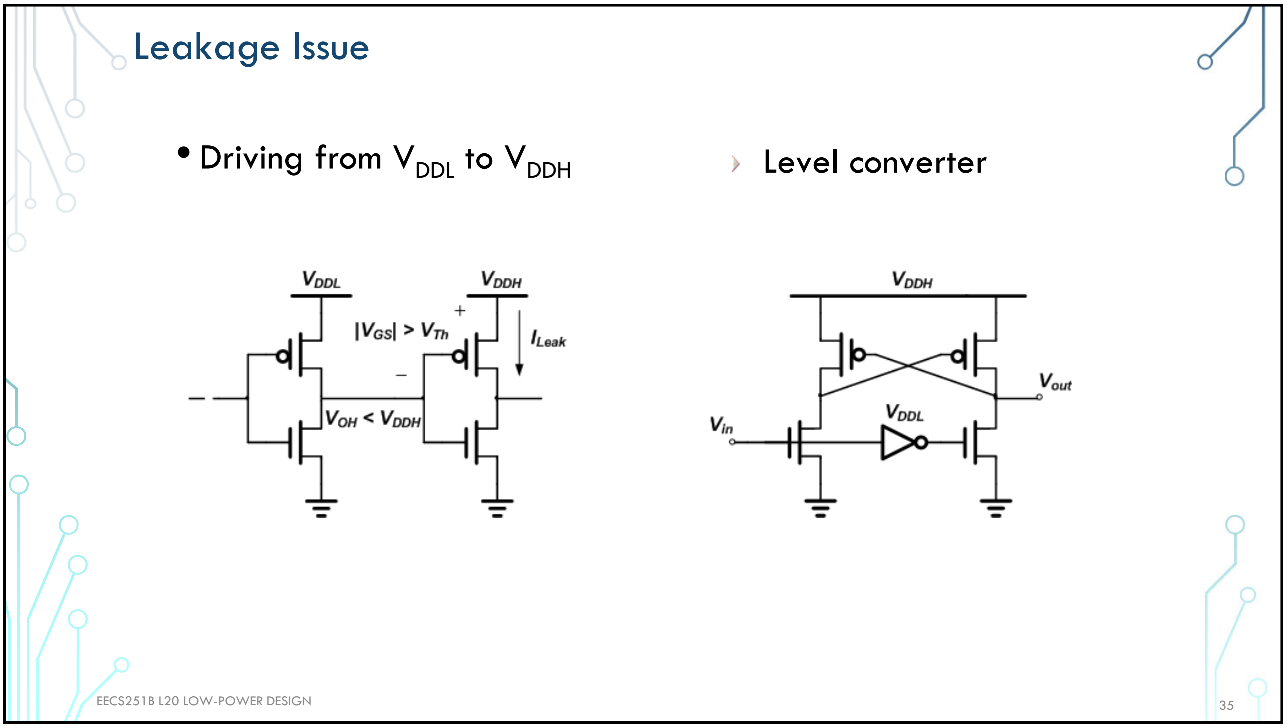

Leakage Issue

- Driving from V_{DDL} to V_{DDH}

› Level converter

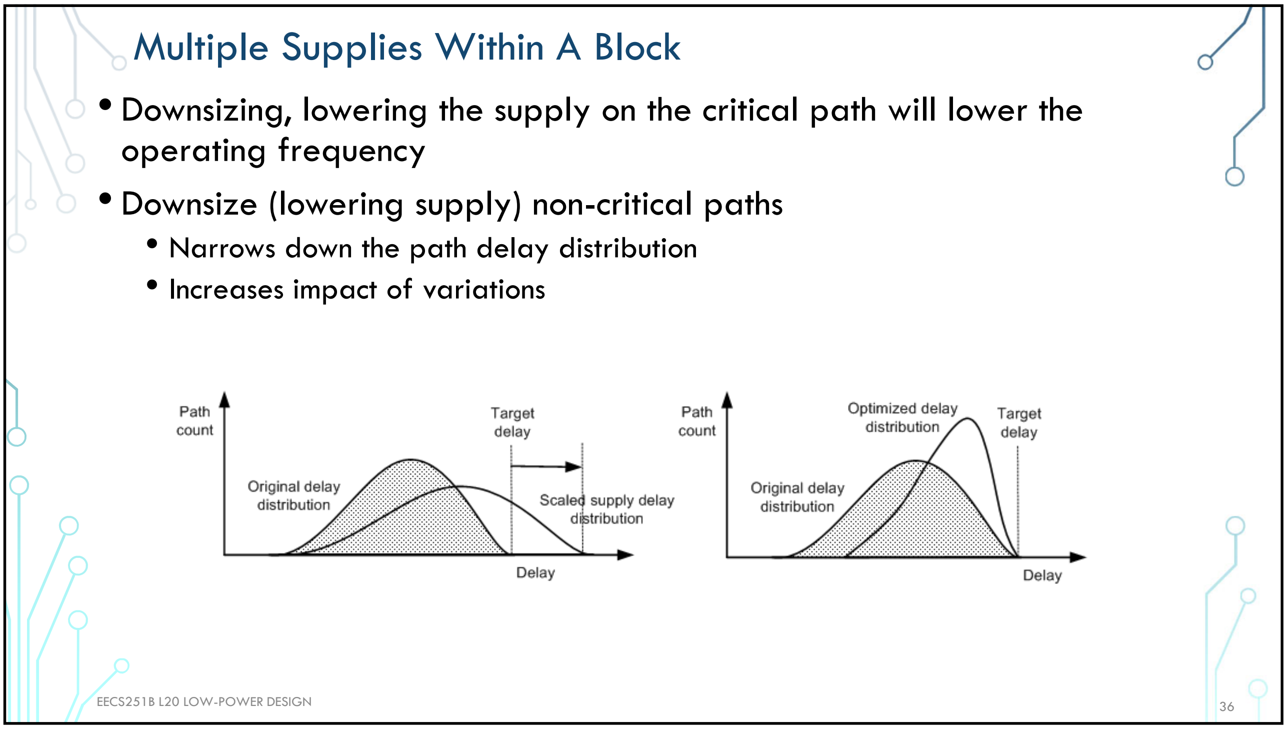

Multiple Supplies Within A Block

- Downsizing, lowering the supply on the critical path will lower the operating frequency
- Downsize (lowering supply) non-critical paths
  - Narrows down the path delay distribution
  - Increases impact of variations

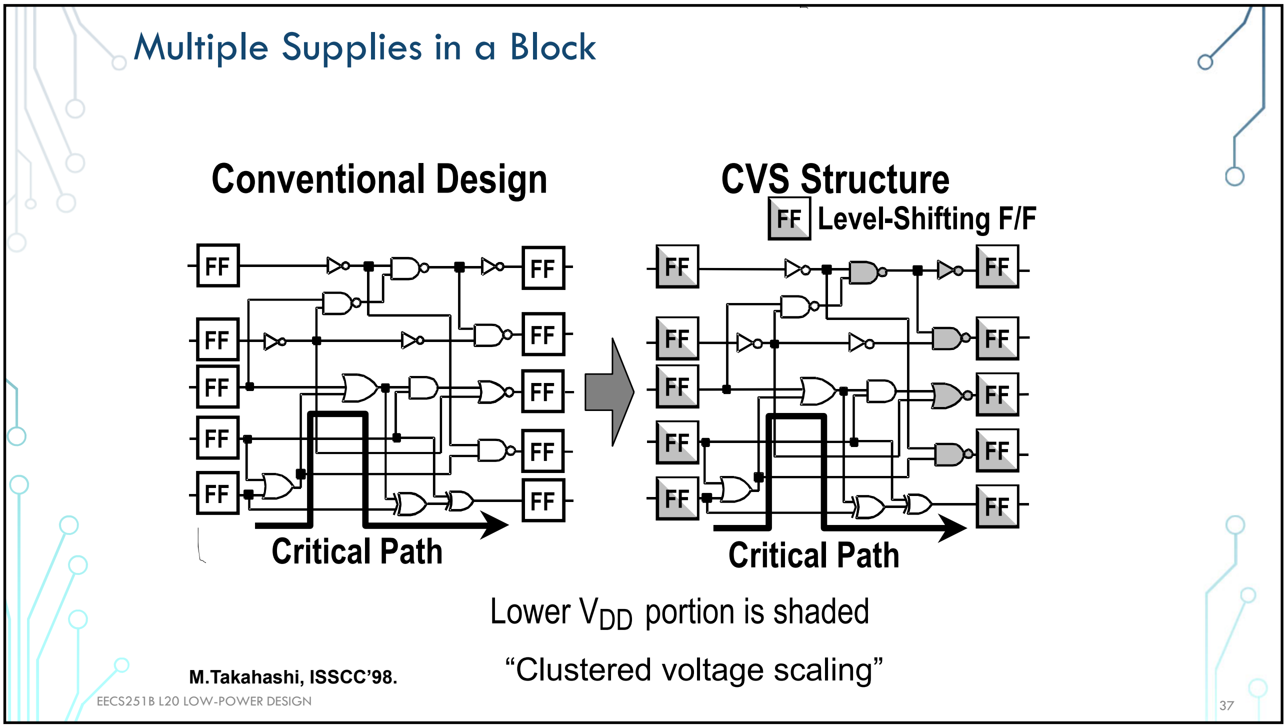

Multiple Supplies in a Block

Lower V_{DD} portion is shaded

“Clustered voltage scaling”

M.Takahashi, ISSCC’98.

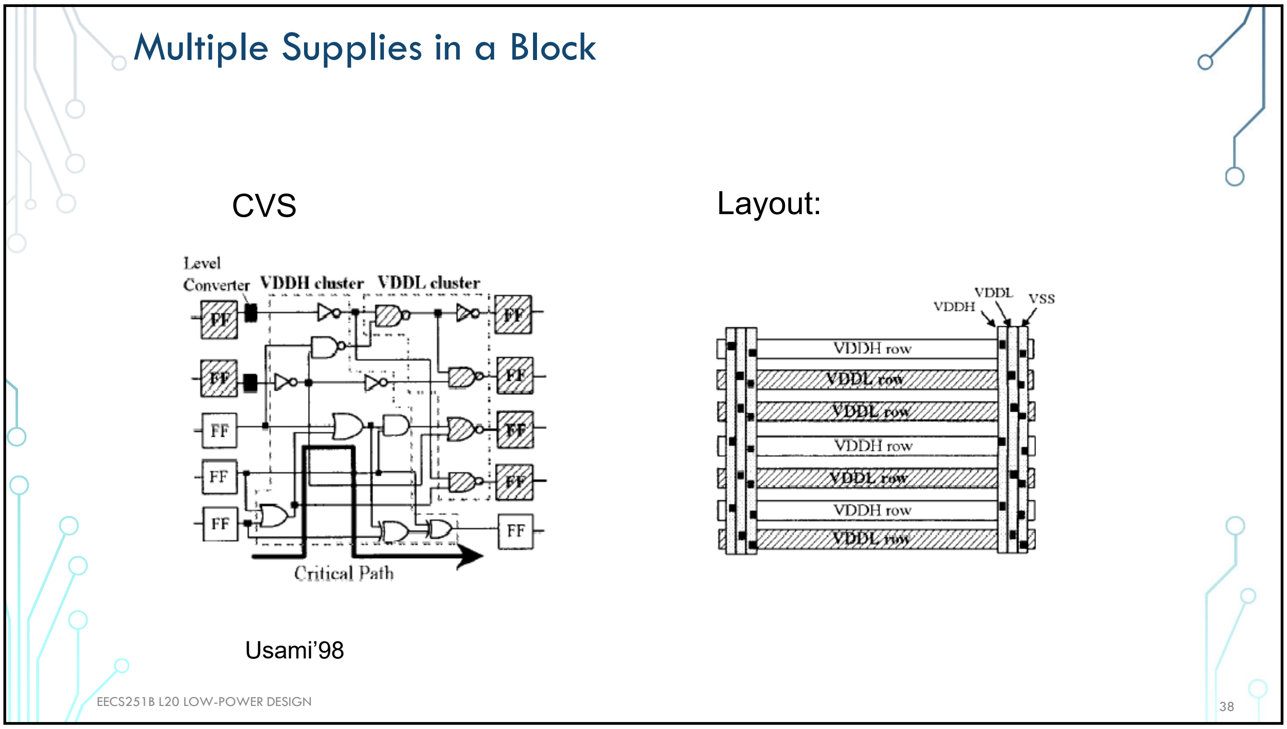

Multiple Supplies in a Block

CVS

Layout:

Usami’98

Level-Converting Flip-Flop

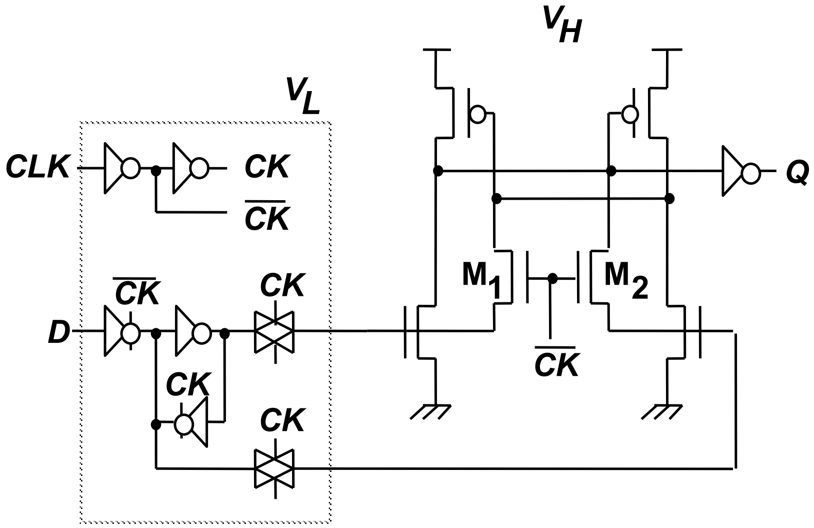

Summary

- Power-performance tradeoffs
  - Sizing
  - Supplies
  - Thresholds
- Lowering supplies
- Multiple supply voltages

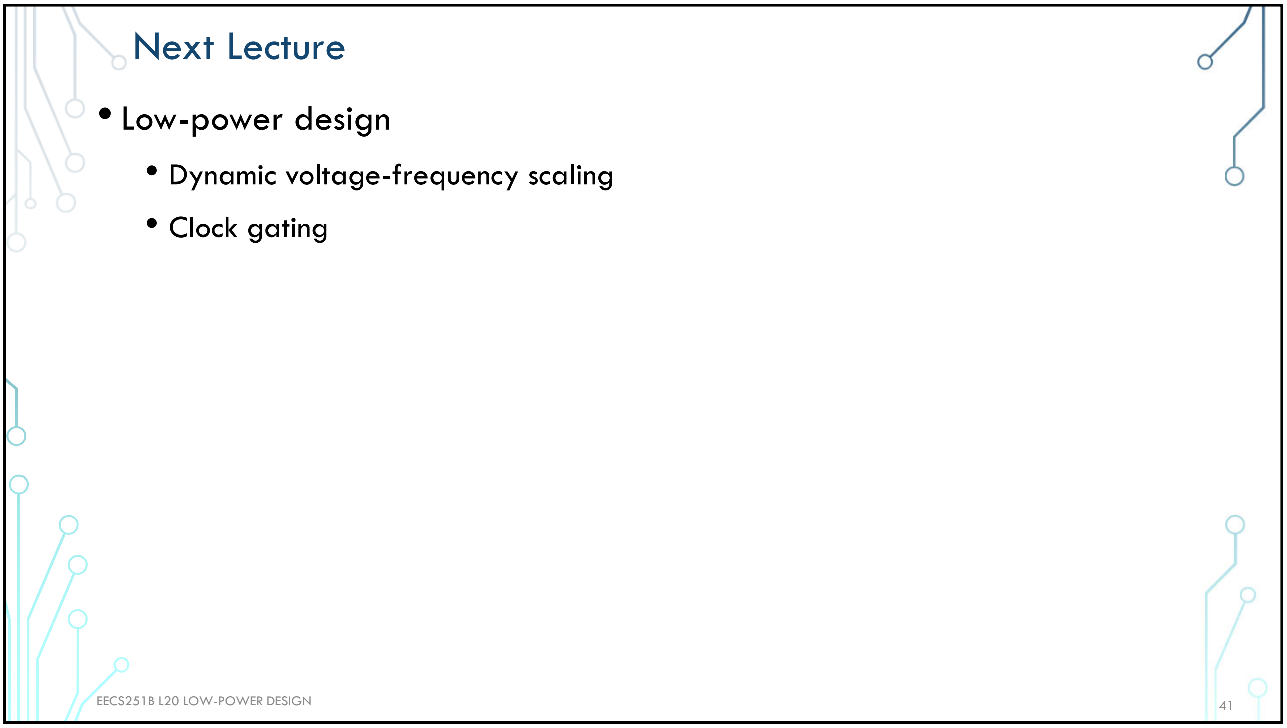

Next Lecture

- Low-power design
  - Dynamic voltage-frequency scaling
  - Clock gating